

Rubric:

Picks a sensible approach for the task	[]
Meets core reporting requirements	[]
Meets core performance requirements [1]	[]
Avoids scope creep	[]
Follow best practices for cross-validation	[]

Pre-trained model used: <https://huggingface.co/sentence-transformers/all-MiniLM-L6-v2>

Final Validation Results:
Loss: 0.04410463944077492, Accuracy: 98%

Test Results:
Loss: 0.11611884087324142, Accuracy: 91%

Configs for both:
learning:
 step_size: 0.07
 batch_size: 500
 num_iters: 300
 decay_rate: 0.999
data:
 noise_stddev: 0.1
mlp:
 num_hidden_layers: 8
 hidden_layer_width: 30
display:
 refresh_rate: 1

Other notes:

- Adam optimizer and binary cross entropy used
- train/val/test split: 10000/2000/7600
- general structure: data -> transformer -> mlp